

▼ Lab5A - Constructing a CNN Network

For spatial data for example image or video data, Convolutional Neural Network (CNN or ConvNet) performs much better than standard neural network. In this practical, we shall learn how to build a CNN Network.

Objectives:

1. Learn how to build a convolutional neural network (CNN)
2. Learn how to build a network or layer using `sequential`

Remember to **enable the GPU** (Edit > Notebook setting > GPU) to ensure short training time.

```
1 from google.colab import drive
2 drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call `drive.mount("/content/gdrive", force_remount=True)`

```
1 cd "/content/gdrive/MyDrive/UCCD3074_Labs/UCCD3074_Lab5"

/content/gdrive/MyDrive/UCCD3074_Labs/UCCD3074_Lab5
```

Import the required libraries

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
```

```
7
8 import torchvision
9 import torchvision.transforms as transforms
10
11 import torch.optim as optim
12 from torch.utils.data import DataLoader
13
14 from torchsummary import summary
15
16 from cifar10 import CIFAR10
17
18 %load_ext autoreload
19 %autoreload 2
```

▼ SECTION 1. DEFINING A CNN MODULE WITH `torch.nn.Module`

In this section, we create a CNN network using `nn.Module`. The `Module` is the main building block, it defines the base class for all neural network and you MUST subclass it.

1.1 Build the network

Exercise. Build the following CNN. You will need to following modules:

- To define a conv2d layer: [`torch.nn.Conv2d\(in_channel, out_channel, kernel_size, stride=1, padding=0\)`](#).
 - `in_channel`: number of channels in the input tensor.
 - `out_channel`: number of channels in the output tensor. This is equivalent to the number of filters in the current convolutional layer.
 - `kernel_size`: size of the filter (`f`).
 - `stride`: stride (`s`). Default value is 1.
 - `padding`: padding (`p`). Default value is 0.
- To define a max pooling layer: [`torch.nn.functional.max_pool2d\(x, kernel_size, stride=None, padding=0\)`](#).

- `x`: input tensor of shape (b, c, h, w) . This is required as this is a functional operation.
- `kernel_size`: size of the filter (f).
- `stride`: stride (s). Default value is `kernel_size`.
- `padding`: padding (p). Default value is 0.
- To define a linear layer: [`torch.nn.Linear \(in_features, out_features\)`](#)
 - `in_features`: size of each input sample. This is equivalent to the number of units or signals in the previous layer.
 - `out_features`: size of the output sample. This is equivalent to the number of units / neurons in the current layer.
- To define the global average pooling: [`torch.mean \(x, dim\)`](#)
 - `x`: the input tensor
 - `dim`: the dimensions to reduce. For the input tensor is (b, c, h, w) , to compute the mean of the spatial dimensions h and w , set `dim = [2, 3]`. This will compute the mean for the spatial dimensions and output a tensor of shape $(b, c, 1, 1)$. Then, use [`torch.squeeze`](#) to remove the two empty dimensions to get a tensor of shape (b, c) .
- Alternatively, the global average pooling can be defined using the following command: [`torch.nn.AdaptiveAvgPool2d \(output_size\)`](#)
 - `output_size`: the target output size (o). The layer will configure the kernel size as $(input_size + target_size - 1) // target_size$ to generate an output tensor of shape `output_size`.

Network Architecture

Layer	Name	Description	OutputShape
-	Input	-	(?, 3, 32, 32)
1	conv1	Conv2d (k=32, f=3, s=1, p=1)	(?, 32, 32, 32)
		relu	(?, 32, 32, 32)
2	conv2	Conv2d (k=32, f=3, s=1, p=1)	(?, 32, 32, 32)
		relu	(?, 32, 32, 32)
	pool1	maxpool (f=2, s=2, p=0)	(?, 32, 16, 16)
3	conv3	Conv2d (k=64, f=3, s=1, p=1)	(?, 64, 16, 16)

Layer	Name	Description	OutputShape
		relu	(?, 64, 16, 16)
4	conv4	Conv2d (k=64, f=3, s=1, p=1)	(?, 64, 16, 16)
		relu	(?, 64, 16, 16)
	global_pool	AdaptiveAvgPool (o=(1,1))	(?, 64, 1, 1)
		view	(?, 64)
5	fc1	Linear (#units=10)	(?, 10)

Notes: k : number of filters, f : filter or kernel size, s : stride, p : padding, o : output shape

```

1 device = "cuda" if torch.cuda.is_available() else "cpu"

1 class Net1(nn.Module):
2     def __init__(self):
3         # call super constructor
4         super().__init__()
5
6         # create the conv1 layer
7         self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1)
8
9         # create the conv2 layer
10        self.conv2 = nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1)
11
12        # create the conv3 layer
13        self.conv3 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
14
15        # create the conv4 layer
16        self.conv4 = nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1)
17
18        # create the global pooling layer
19        self.global_pool = nn.AdaptiveAvgPool2d((1,1))
20
21        # fully connected layer
22        self.fc1 = nn.Linear(64, 10)
23

```

```
24     def forward(self, x):
25
26         # conv1 layer
27         x = F.relu(self.conv1(x))
28
29         # conv2 layer
30         x = F.relu(self.conv2(x))
31
32         # pooling layer
33         x = F.max_pool2d(x, 2, 2)
34
35         # conv3 layer
36         x = F.relu(self.conv3(x))
37
38         # conv4 layer
39         x = F.relu(self.conv4(x))
40
41         # global pooling
42         # x = torch.mean(x, [2, 3])
43         x = self.global_pool(x)
44
45         # remove the spatial dimension
46         x = x.squeeze()
47
48         # fc1 layer
49         x = self.fc1(x)
50
51         return x
```

Create the network and test it

```
1 net1 = Net1()
2 out = net1(torch.randn(4, 3, 32, 32))
```

Display the network

```
1 print(net1)
```

```
Net1(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv4): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (global_pool): AdaptiveAvgPool2d(output_size=(1, 1))
  (fc1): Linear(in_features=64, out_features=10, bias=True)
)
```

```
1 summary(net1, input_size=(3, 32, 32), device="cpu")
```

```
-----
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	896
Conv2d-2	[-1, 32, 32, 32]	9,248
Conv2d-3	[-1, 64, 16, 16]	18,496
Conv2d-4	[-1, 64, 16, 16]	36,928
AdaptiveAvgPool2d-5	[-1, 64, 1, 1]	0
Linear-6	[-1, 10]	650

```
=====
Total params: 66,218
Trainable params: 66,218
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 0.75
Params size (MB): 0.25
Estimated Total Size (MB): 1.01
-----
```

▼ 1.2 Load the dataset

Exercise. Load the dataset

1. Load the dataset. Define the following transformation pipeline to
 - Convert an image (numpy array with range (0, 255)) to a tensor, and
 - Normalize the tensor with mean = 0.5 and std = 1.0

```
1 # transform the model
2 transform = transforms.Compose([
3     transforms.ToTensor(),
4     transforms.Normalize((0.5, 0.5, 0.5), (1., 1., 1.))
5 ])
6
7 # Load the dataset
8 trainset = CIFAR10(train=True, transform=transform, download=True, num_samples=10000)
9 testset = CIFAR10(train=False, transform=transform, download=True, num_samples=2000)
10
11 print('Size of trainset:', len(trainset))
12 print('Size of testset:', len(testset))
```

```
Files already downloaded and verified
Files already downloaded and verified
Size of trainset: 10000
Size of testset: 2000
```

2. Create the dataloader for train set and test set. Use a batch size of 16, enable shuffle, apply the transformation pipeline defined above and use 2 cpu workers to load the datasets.

```
1 trainloader = DataLoader(trainset, batch_size=16, shuffle=True, num_workers=2)
2 testloader = DataLoader(testset, batch_size=16, shuffle=True, num_workers=2)
```

▼ 1.3 Train the model

Exercise. Complete the training function below

```
1 def train(net, trainloader, num_epochs=15, lr=0.1, momentum=0.9):
2
3     loss_iterations = int(np.ceil(len(trainloader)/3))
4
5     # transfer model to GPU
6     net = net.to(device)
7
8     # set the optimizer. Use the SGD optimizer. Use the lr and momentum settings passed by the user
9     optimizer = optim.SGD(net.parameters(), lr=lr, momentum=momentum)
10
11     # set to training mode
12     net.train()
13
14     # variables
15     best_loss = np.inf
16     saturate_count = 0
17
18     # train the network
19     for e in range(num_epochs):
20
21         running_loss = 0
22         running_count = 0
23
24         # for all batch samples
25         for i, (inputs, labels) in enumerate(trainloader):
26
27             # Clear all the gradient to zero
28             optimizer.zero_grad()
29
30             # transfer data to GPU
31             inputs = inputs.to(device)
32             labels = labels.to(device)
33
```



```

34         # forward propagation to get h
35         outs = net(inputs)
36
37         # compute loss
38         loss = F.cross_entropy(outs, labels)
39
40         # backpropagation to get gradients of all parameters
41         loss.backward()
42
43         # update parameters
44         optimizer.step()
45
46         # get the loss
47         running_loss += loss.item()
48         running_count += 1
49
50         # display the averaged loss value
51         if i % loss_iterations == loss_iterations-1 or i == len(trainloader) - 1:
52             train_loss = running_loss / running_count
53             running_loss = 0.
54             running_count = 0.
55             print(f'[Epoch {e+1:2d} Iter {i+1:5d}/{len(trainloader)}]: train_loss = {train_loss:.4f}')
56
57     print("Training completed.")

```

Now, train the model with a maximum number of epochs of 50. The training will stop once it converge and may stop earlier. Use a learning rate of 0.01 and momentum of 0.9.

```
1 train(net1, trainloader, num_epochs=15, lr=0.01, momentum=0.9)
```

```

[Epoch 1 Iter 209/625]: train_loss = 2.3033
[Epoch 1 Iter 418/625]: train_loss = 2.2374
[Epoch 1 Iter 625/625]: train_loss = 2.1177
[Epoch 2 Iter 209/625]: train_loss = 2.0776
[Epoch 2 Iter 418/625]: train_loss = 2.0541
[Epoch 2 Iter 625/625]: train_loss = 1.9898

```

```
[Epoch 3 Iter 209/625]: train_loss = 1.9296
[Epoch 3 Iter 418/625]: train_loss = 1.8712
[Epoch 3 Iter 625/625]: train_loss = 1.8263
[Epoch 4 Iter 209/625]: train_loss = 1.8003
[Epoch 4 Iter 418/625]: train_loss = 1.7818
[Epoch 4 Iter 625/625]: train_loss = 1.7367
[Epoch 5 Iter 209/625]: train_loss = 1.7340
[Epoch 5 Iter 418/625]: train_loss = 1.7119
[Epoch 5 Iter 625/625]: train_loss = 1.6874
[Epoch 6 Iter 209/625]: train_loss = 1.6954
[Epoch 6 Iter 418/625]: train_loss = 1.6175
[Epoch 6 Iter 625/625]: train_loss = 1.6409
[Epoch 7 Iter 209/625]: train_loss = 1.6300
[Epoch 7 Iter 418/625]: train_loss = 1.5704
[Epoch 7 Iter 625/625]: train_loss = 1.5864
[Epoch 8 Iter 209/625]: train_loss = 1.5217
[Epoch 8 Iter 418/625]: train_loss = 1.5524
[Epoch 8 Iter 625/625]: train_loss = 1.5273
[Epoch 9 Iter 209/625]: train_loss = 1.5390
[Epoch 9 Iter 418/625]: train_loss = 1.4959
[Epoch 9 Iter 625/625]: train_loss = 1.4553
[Epoch 10 Iter 209/625]: train_loss = 1.4317
[Epoch 10 Iter 418/625]: train_loss = 1.4425
[Epoch 10 Iter 625/625]: train_loss = 1.3974
[Epoch 11 Iter 209/625]: train_loss = 1.3811
[Epoch 11 Iter 418/625]: train_loss = 1.3577
[Epoch 11 Iter 625/625]: train_loss = 1.3808
[Epoch 12 Iter 209/625]: train_loss = 1.3373
[Epoch 12 Iter 418/625]: train_loss = 1.3626
[Epoch 12 Iter 625/625]: train_loss = 1.3055
[Epoch 13 Iter 209/625]: train_loss = 1.3144
[Epoch 13 Iter 418/625]: train_loss = 1.2533
[Epoch 13 Iter 625/625]: train_loss = 1.2920
[Epoch 14 Iter 209/625]: train_loss = 1.2453
[Epoch 14 Iter 418/625]: train_loss = 1.2586
[Epoch 14 Iter 625/625]: train_loss = 1.2454
[Epoch 15 Iter 209/625]: train_loss = 1.2000
[Epoch 15 Iter 418/625]: train_loss = 1.1960
[Epoch 15 Iter 625/625]: train_loss = 1.2186
Training completed.
```

▼ 3. Evaluate the model

Now let's evaluate the model. Remember that a 2-layered neural network only achieves an accuracy of around 38%. With a CNN architecture, you should be able to achieve a higher accuracy of more than 50%.

Exercise. Complete the function `evaluate` below to evaluate the model on the test loader

```
1 def evaluate(net, testloader):
2
3     # set to evaluation mode
4     net.eval()
5
6     # running_correct
7     running_corrects = 0
8
9     # Repeat for all batch data in the test set
10    for inputs, targets in testloader:
11
12        # transfer to the GPU
13        inputs = inputs.to(device)
14        targets = targets.to(device)
15
16        # # disable gradient computation
17        with torch.no_grad():
18
19            # perform inference
20            outputs = net(inputs)
21
22            # predict as the best result
23            _, predicted = torch.max(outputs, 1)
24
25            running_corrects += (targets == predicted).double().sum()
26
27
28    print('Accuracy = {:.2f}%'.format(100*running_corrects/len(testloader.dataset)))
```

Now, let's evaluate our model.

```
1 evaluate(net1, testloader)
```

```
Accuracy = 52.65%
```

➤ 2. CREATING A CNN NETWORK DIRECTLY USING `torch.nn.Sequential`

In this section, we shall learn how to create a network using `torch.nn.Sequential`. `Sequential` is a container of `Modules` that can be stacked together and run at the same time.

```
net = nn.Sequential(  
    nn.Conv2d(...),  
    nn.ReLU(),  
    ...  
)  
  
x = ... # get the input tensor  
output = net(x) # perform inference
```

We can see immediately that it is a very convenient way to build a network.

Limitations: Note that you cannot add functional operations (e.g., `torch.relu`) into a `Sequential` model. If the `nn` module version does not exist for the function, then you have to create your own `nn` module for the function.

Exercise. Reimplement the network above using `torch.nn.Sequential`.

Since you cannot use functional operations for `Sequential` models, you use their corresponding module versions:

- `torch.nn.functional.max_pool2d` --> [torch.nn.MaxPool2d](#)
- `torch.nn.functional.relu` --> [torch.nn.ReLU](#)
- `torch.view` --> [nn.Flatten](#)

```

1 net2 = nn.Sequential(
2
3     nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1), # conv1
4     nn.ReLU(), # relu
5     nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1), # conv2
6     nn.ReLU(), # relu
7
8     nn.MaxPool2d(kernel_size=2, stride=2, padding=0),      # max pool
9
10    nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1), # conv3
11    nn.ReLU(), # relu
12    nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1), # conv4
13    nn.ReLU(), # relu
14
15    nn.AdaptiveAvgPool2d((1,1)), # global average pooling
16    nn.Flatten(), # flatten
17
18    nn.Linear(64, 10) # fc1
19 )

```

```
1 summary(net2, (3, 32, 32), device = "cpu")
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	896
ReLU-2	[-1, 32, 32, 32]	0
Conv2d-3	[-1, 32, 32, 32]	9,248
ReLU-4	[-1, 32, 32, 32]	0
MaxPool2d-5	[-1, 32, 16, 16]	0
Conv2d-6	[-1, 64, 16, 16]	18,496
ReLU-7	[-1, 64, 16, 16]	0

Conv2d-8	[-1, 64, 16, 16]	36,928
ReLU-9	[-1, 64, 16, 16]	0
AdaptiveAvgPool2d-10	[-1, 64, 1, 1]	0
Flatten-11	[-1, 64]	0
Linear-12	[-1, 10]	650

```

=====
Total params: 66,218
Trainable params: 66,218
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 1.56
Params size (MB): 0.25
Estimated Total Size (MB): 1.83
-----

```

▼ Train the model

```
1 train(net2, trainloader, num_epochs=15, lr=0.01, momentum=0.9)
```

```

[Epoch 1 Iter 209/625]: train_loss = 2.3037
[Epoch 1 Iter 418/625]: train_loss = 2.2817
[Epoch 1 Iter 625/625]: train_loss = 2.1324
[Epoch 2 Iter 209/625]: train_loss = 2.1070
[Epoch 2 Iter 418/625]: train_loss = 2.0628
[Epoch 2 Iter 625/625]: train_loss = 2.0199
[Epoch 3 Iter 209/625]: train_loss = 1.9433
[Epoch 3 Iter 418/625]: train_loss = 1.8955
[Epoch 3 Iter 625/625]: train_loss = 1.8435
[Epoch 4 Iter 209/625]: train_loss = 1.8034
[Epoch 4 Iter 418/625]: train_loss = 1.7735
[Epoch 4 Iter 625/625]: train_loss = 1.7826
[Epoch 5 Iter 209/625]: train_loss = 1.7399
[Epoch 5 Iter 418/625]: train_loss = 1.7113
[Epoch 5 Iter 625/625]: train_loss = 1.7396
[Epoch 6 Iter 209/625]: train_loss = 1.6783
[Epoch 6 Iter 418/625]: train_loss = 1.6685
[Epoch 6 Iter 625/625]: train_loss = 1.6208

```

```
[Epoch 7 Iter 209/625]: train_loss = 1.6188
[Epoch 7 Iter 418/625]: train_loss = 1.6320
[Epoch 7 Iter 625/625]: train_loss = 1.5998
[Epoch 8 Iter 209/625]: train_loss = 1.5504
[Epoch 8 Iter 418/625]: train_loss = 1.5874
[Epoch 8 Iter 625/625]: train_loss = 1.5513
[Epoch 9 Iter 209/625]: train_loss = 1.5319
[Epoch 9 Iter 418/625]: train_loss = 1.4825
[Epoch 9 Iter 625/625]: train_loss = 1.5077
[Epoch 10 Iter 209/625]: train_loss = 1.4713
[Epoch 10 Iter 418/625]: train_loss = 1.4456
[Epoch 10 Iter 625/625]: train_loss = 1.4627
[Epoch 11 Iter 209/625]: train_loss = 1.4206
[Epoch 11 Iter 418/625]: train_loss = 1.4193
[Epoch 11 Iter 625/625]: train_loss = 1.3976
[Epoch 12 Iter 209/625]: train_loss = 1.3967
[Epoch 12 Iter 418/625]: train_loss = 1.3494
[Epoch 12 Iter 625/625]: train_loss = 1.3385
[Epoch 13 Iter 209/625]: train_loss = 1.3304
[Epoch 13 Iter 418/625]: train_loss = 1.3122
[Epoch 13 Iter 625/625]: train_loss = 1.3069
[Epoch 14 Iter 209/625]: train_loss = 1.2950
[Epoch 14 Iter 418/625]: train_loss = 1.2506
[Epoch 14 Iter 625/625]: train_loss = 1.2692
[Epoch 15 Iter 209/625]: train_loss = 1.2458
[Epoch 15 Iter 418/625]: train_loss = 1.2469
[Epoch 15 Iter 625/625]: train_loss = 1.2405
Training completed.
```

▼ Evaluate the model on the test set

```
1 evaluate(net2, testloader)
```

Accuracy = 55.20%

3. Using a function to create a customizable BLOCK module

In the following, we group the convolutional layers in the network above into 2 blocks:

- conv1 and conv2 --> block_1
- conv3 and conv4 --> block_2.

This is how the network looks:

NETWORK (with block)				
Block	Layer	Name	Description	OutputShape
input	-	-	-	(?, 3, 32, 32)
block_1 (blk_cin=3, blk_cout=32)	1	conv1	Conv2d (in_channels=3, out_channels=32,f=3,s=1,p=1)	(?, 32, 32, 32)
	-	ReLU	relu	(?, 32, 32, 32)
	2	conv2	Conv2d (in_channels=32, out_channels=32,f=3,s=1,p=1)	(?, 32, 32, 32)
	-	ReLU	relu	(?, 32, 32, 32)
-	-	pool1	maxpool (f=2,s=2,p=0)	(?, 32, 16, 16)
block_2 (blk_cin=32, blk_cout=64)	3	conv1	Conv2d (in_channels=32, out_channels=64,f=3,s=1,p=1)	(?, 64, 16, 16)
	-	ReLU	relu	(?, 64, 16, 16)
	4	conv2	Conv2d (in_channels=64, out_channels=64, f=3,s=1,p=1)	(?, 64, 16, 16)
	-	ReLU	relu	(?, 64, 16, 16)
	-	global_pool	AdaptiveAvgPool, o=(1,1)	(?, 64, 1, 1)
	-	-	view	(?, 64)
	5	fc1	Linear(#units=10)	(?, 10)
	-	-	view	(?, 10)

Notes: *k*: number of filters, *f*: filter or kernel size, *s*: stride, *p*: padding, *o*: output shape

Comparing block_1 and block_2, we find that they have similar structure


```
conv (blk_cin, blk_cout) --> relu --> conv (blk_cout, blk_cout) --> relu
```

where

- blk_cin=3 and blk_cout=32 for block1
- blk_cin=32 and blk_cout=64 for block2

▼ The BUILD_BLOCK function

Rather than declaring each layer individually, you create a function `build_block` to create a block of layers. The function receives `in_ch`, the number of channels in the input tensor and `out_ch`, the number of channels in the output tensor. In the following, complete the `build_block` function by returning a `nn.Sequential` module that builds and returns the following block

BLOCK (blk_cin, blk_cout)
Conv2d (in_channels=blk_cin, output_channels=blk_cout, f=3, s=1,p=1)
ReLU ()
Conv2d (in_channels=blk_cout, output_channels=blk_cout, f=3, s=1,p=1)
ReLU ()

```
1 def build_block(blk_cin, blk_cout):
2     block = nn.Sequential(
3         nn.Conv2d(blk_cin, blk_cout, kernel_size=3, stride=1, padding=1),
4         nn.ReLU(),
5         nn.Conv2d(blk_cout, blk_cout, kernel_size=3, stride=1, padding=1),
6         nn.ReLU(),
7     )
8     return block
```

```
1 block1 = build_block(3, 32)
2 print(block1)
```

```
Sequential(
  (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
  (2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
)
```

```
(3): ReLU()
)
```

▼ Constructing the network with BUILD_BLOCK

Now, construct NETWORK (with block). Use the build_block function to construct block1 and block2

```
1 class Net3(nn.Module):
2     def __init__(self):
3         super().__init__()
4
5         # define block 1
6         self.conv_block1 = build_block(3, 32)
7
8         # define block 2
9         self.conv_block2 = build_block(32, 64)
10
11        # define global_pool
12        self.global_pool = nn.AdaptiveAvgPool2d((1,1))
13
14        # define fc1
15        self.fc1 = nn.Linear(64, 10)
16
17    def forward(self, x):
18        # block 1
19        x = self.conv_block1(x)
20
21        # max pool
22        x = F.max_pool2d(x, kernel_size=2, stride=2, padding=0)
23
24        # block 2
25        x = self.conv_block2(x)
26
27        # global pool
28        x = self.global_pool(x)
```

```

29
30     # view
31     x = x.view(x.size(0), -1)
32
33     # fc1
34     x = self.fc1(x)
35
36     return x

```

```

1 net3 = Net3()
2 summary(net3, (3, 32, 32), device="cpu")

```

```

-----
      Layer (type)          Output Shape          Param #
=====
      Conv2d-1             [-1, 32, 32, 32]           896
      ReLU-2               [-1, 32, 32, 32]            0
      Conv2d-3             [-1, 32, 32, 32]          9,248
      ReLU-4               [-1, 32, 32, 32]            0
      Conv2d-5             [-1, 64, 16, 16]         18,496
      ReLU-6               [-1, 64, 16, 16]            0
      Conv2d-7             [-1, 64, 16, 16]         36,928
      ReLU-8               [-1, 64, 16, 16]            0
      AdaptiveAvgPool2d-9  [-1, 64, 1, 1]             0
      Linear-10            [-1, 10]                   650
=====
Total params: 66,218
Trainable params: 66,218
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 1.50
Params size (MB): 0.25
Estimated Total Size (MB): 1.76
-----

```

▼ Train the model

```
1 train(net3, trainloader, num_epochs=15, lr=0.01, momentum=0.9)
```

```
[Epoch 1 Iter 209/625]: train_loss = 2.3043
[Epoch 1 Iter 418/625]: train_loss = 2.2616
[Epoch 1 Iter 625/625]: train_loss = 2.1163
[Epoch 2 Iter 209/625]: train_loss = 2.0972
[Epoch 2 Iter 418/625]: train_loss = 2.0199
[Epoch 2 Iter 625/625]: train_loss = 1.9518
[Epoch 3 Iter 209/625]: train_loss = 1.9106
[Epoch 3 Iter 418/625]: train_loss = 1.8369
[Epoch 3 Iter 625/625]: train_loss = 1.8175
[Epoch 4 Iter 209/625]: train_loss = 1.8142
[Epoch 4 Iter 418/625]: train_loss = 1.7458
[Epoch 4 Iter 625/625]: train_loss = 1.7503
[Epoch 5 Iter 209/625]: train_loss = 1.7100
[Epoch 5 Iter 418/625]: train_loss = 1.6953
[Epoch 5 Iter 625/625]: train_loss = 1.6882
[Epoch 6 Iter 209/625]: train_loss = 1.6642
[Epoch 6 Iter 418/625]: train_loss = 1.6411
[Epoch 6 Iter 625/625]: train_loss = 1.6486
[Epoch 7 Iter 209/625]: train_loss = 1.6068
[Epoch 7 Iter 418/625]: train_loss = 1.6026
[Epoch 7 Iter 625/625]: train_loss = 1.5702
[Epoch 8 Iter 209/625]: train_loss = 1.5390
[Epoch 8 Iter 418/625]: train_loss = 1.5445
[Epoch 8 Iter 625/625]: train_loss = 1.5442
[Epoch 9 Iter 209/625]: train_loss = 1.4911
[Epoch 9 Iter 418/625]: train_loss = 1.5032
[Epoch 9 Iter 625/625]: train_loss = 1.4851
[Epoch 10 Iter 209/625]: train_loss = 1.4244
[Epoch 10 Iter 418/625]: train_loss = 1.4538
[Epoch 10 Iter 625/625]: train_loss = 1.4037
[Epoch 11 Iter 209/625]: train_loss = 1.3941
[Epoch 11 Iter 418/625]: train_loss = 1.3626
[Epoch 11 Iter 625/625]: train_loss = 1.3757
[Epoch 12 Iter 209/625]: train_loss = 1.3141
[Epoch 12 Iter 418/625]: train_loss = 1.3364
[Epoch 12 Iter 625/625]: train_loss = 1.3011
[Epoch 13 Iter 209/625]: train_loss = 1.2824
```

```
[Epoch 13 Iter 418/625]: train_loss = 1.2968
[Epoch 13 Iter 625/625]: train_loss = 1.2925
[Epoch 14 Iter 209/625]: train_loss = 1.2376
[Epoch 14 Iter 418/625]: train_loss = 1.1726
[Epoch 14 Iter 625/625]: train_loss = 1.2761
[Epoch 15 Iter 209/625]: train_loss = 1.1994
[Epoch 15 Iter 418/625]: train_loss = 1.2135
[Epoch 15 Iter 625/625]: train_loss = 1.1916
Training completed.
```

▼ Evaluate the model

```
1 evaluate(net3, testloader)
```

```
Accuracy = 52.65%
```

▼ 4. Create a customizable BLOCK module

We can also build a block by constructing a module called `BLOCK` where we can specify the `blk_cin` and `blk_cout` when constructing the block.

The BLOCK module

Implement the BLOCK module.

BLOCK (blk_cin, blk_cout)

```
Conv2d (in_channels=blk_cin, output_channels=blk_cout, f=3, s=1,p=1)
```

```
ReLU ()
```

```
Conv2d (in_channels=blk_cout, output_channels=blk_cout, f=3, s=1,p=1)
```

```
ReLU ()
```

```
1 class BLOCK(nn.Module):
```

```

2     def __init__(self, blk_cin, blk_cout):
3
4         super().__init__()
5
6         self.conv1 = nn.Conv2d(blk_cin, blk_cout, kernel_size=3, stride=1, padding=1)
7         self.conv2 = nn.Conv2d(blk_cout, blk_cout, kernel_size=3, stride=1, padding=1)
8
9     def forward(self, x):
10
11         x = F.relu(self.conv1(x))
12         x = F.relu(self.conv2(x))
13
14         return x

```

▼ Construcing network with the BLOCK module

Now, let's build the network using the `BLOCK` module that we have just created.

```

1 class Net4(nn.Module):
2     def __init__(self):
3         super().__init__()
4
5         # define block 1
6         self.conv_block1 = BLOCK(3, 32)
7
8         # define block 2
9         self.conv_block2 = BLOCK(32, 64)
10
11        # define global_pool
12        self.global_pool = nn.AdaptiveAvgPool2d((1,1))
13
14        # define fc1
15        self.fc1 = nn.Linear(64, 10)
16
17    def forward(self, x):

```

```

18         # block 1
19         x = self.conv_block1(x)
20
21         # max pool
22         x = F.max_pool2d(x, kernel_size=2, stride=2, padding=0)
23
24         # block 2
25         x = self.conv_block2(x)
26
27         # global pool
28         x = self.global_pool(x)
29
30         # view
31         x = x.view(x.size(0), -1)
32
33         # fc1
34         x = self.fc1(x)
35
36         return x

```

```
1 net4 = Net4()
```

▼ Train the model

```
1 train(net4, trainloader, num_epochs=15, lr=0.01, momentum=0.9)
```

```

↳ [Epoch 1 Iter 209/625]: train_loss = 2.3015
   [Epoch 1 Iter 418/625]: train_loss = 2.2380
   [Epoch 1 Iter 625/625]: train_loss = 2.1052
   [Epoch 2 Iter 209/625]: train_loss = 2.0804
   [Epoch 2 Iter 418/625]: train_loss = 2.0293
   [Epoch 2 Iter 625/625]: train_loss = 1.9657
   [Epoch 3 Iter 209/625]: train_loss = 1.9291
   [Epoch 3 Iter 418/625]: train_loss = 1.8748
   [Epoch 3 Iter 625/625]: train_loss = 1.8325

```

```
[Epoch 4 Iter 209/625]: train_loss = 1.7820
[Epoch 4 Iter 418/625]: train_loss = 1.7922
[Epoch 4 Iter 625/625]: train_loss = 1.7768
[Epoch 5 Iter 209/625]: train_loss = 1.7579
[Epoch 5 Iter 418/625]: train_loss = 1.7392
[Epoch 5 Iter 625/625]: train_loss = 1.7319
[Epoch 6 Iter 209/625]: train_loss = 1.6866
[Epoch 6 Iter 418/625]: train_loss = 1.6857
[Epoch 6 Iter 625/625]: train_loss = 1.6418
[Epoch 7 Iter 209/625]: train_loss = 1.6643
[Epoch 7 Iter 418/625]: train_loss = 1.6227
[Epoch 7 Iter 625/625]: train_loss = 1.6382
[Epoch 8 Iter 209/625]: train_loss = 1.6152
[Epoch 8 Iter 418/625]: train_loss = 1.5588
[Epoch 8 Iter 625/625]: train_loss = 1.5742
[Epoch 9 Iter 209/625]: train_loss = 1.5690
[Epoch 9 Iter 418/625]: train_loss = 1.5363
[Epoch 9 Iter 625/625]: train_loss = 1.5349
[Epoch 10 Iter 209/625]: train_loss = 1.4852
[Epoch 10 Iter 418/625]: train_loss = 1.4738
[Epoch 10 Iter 625/625]: train_loss = 1.4711
[Epoch 11 Iter 209/625]: train_loss = 1.4782
[Epoch 11 Iter 418/625]: train_loss = 1.4293
[Epoch 11 Iter 625/625]: train_loss = 1.3731
[Epoch 12 Iter 209/625]: train_loss = 1.4311
[Epoch 12 Iter 418/625]: train_loss = 1.4043
[Epoch 12 Iter 625/625]: train_loss = 1.3351
[Epoch 13 Iter 209/625]: train_loss = 1.3347
[Epoch 13 Iter 418/625]: train_loss = 1.3368
[Epoch 13 Iter 625/625]: train_loss = 1.2996
[Epoch 14 Iter 209/625]: train_loss = 1.2829
[Epoch 14 Iter 418/625]: train_loss = 1.2901
[Epoch 14 Iter 625/625]: train_loss = 1.2729
[Epoch 15 Iter 209/625]: train_loss = 1.2537
[Epoch 15 Iter 418/625]: train_loss = 1.2475
[Epoch 15 Iter 625/625]: train_loss = 1.2435
Training completed.
```

▼ Evaluate the model


```
1 evaluate(net4, testloader)
```

```
Accuracy = 55.00%
```

--- End of Practical ---

✓ 0s completed at 6:07 PM

● ✕